

# **Introduction of Feature Hashing**

**Create a Model Matrix via Feature Hashing with a Formula Interface**


Wush Wu

Taiwan R User Group

# Wush Chi-Hsuan Wu

- Ph.D. Student
  - Display Advertising
  - Large Scale Machine Learning
- Familiar tools: R, C++, [Apache Spark](#)

# Taiwan, a Mountainous Island



source: <http://www.nbcnews.com>

- Highest Mountain: Yushan (3952m)



# Taipei City and Taipei 101



source: [https://c2.staticflickr.com/6/5208/5231215951\\_c0e0036b17\\_b.jpg](https://c2.staticflickr.com/6/5208/5231215951_c0e0036b17_b.jpg)




# New Landmark of Taipei



# Taiwan R User Group

<http://www.meetup.com/Taiwan-R>



# An Example of FeatureHashing

# Sentiment Analysis

- Provided by [Lewis Crouch](#)
- Show the pros of using `FeatureHashing`



# Dataset: IMDB

- Predict the rating of the movie according to the text in review.
- Training Dataset: 25000 reviews
- Binary response: positive or negative.

```
## [1] 1
```

- Cleaned review

```
## [1] "kurosawa is a proved humanitarian this movie is totally about people living in"  
## [2] "poverty you will see nothing but angry in this movie it makes you feel bad but"  
## [3] "still worth all those who s too comfortable with materialization should spend 2"  
## [4] "5 hours with this movie"
```

# Word Segmentation

##	[1]	" "	"kurosawa"	"is"
##	[4]	"a"	"proved"	"humanitarian"
##	[7]	" "	"this"	"movie"
##	[10]	"is"	"totally"	"about"
##	[13]	"people"	"living"	"in"
##	[16]	"poverty"	" "	"you"
##	[19]	"will"	"see"	"nothing"
##	[22]	"but"	"angry"	"in"
##	[25]	"this"	"movie"	" "
##	[28]	"it"	"makes"	"you"
##	[31]	"feel"	"bad"	"but"
##	[34]	"still"	"worth"	" "
##	[37]	"all"	"those"	"who"
##	[40]	"s"	"too"	"comfortable"
##	[43]	"with"	"materialization"	"should"
##	[46]	"spend"	"2"	"5"
##	[49]	"hours"	"with"	"this"
##	[52]	"movie"	" "	

# Word Segmentation and Feature Extraction

ID	MESSAGE	HATE	BAD	LIKE	GRATUITOUS
"5814_8"	TRUE	TRUE	TRUE	TRUE	FALSE
"2381_9"	FALSE	FALSE	FALSE	TRUE	FALSE
"7759_3"	FALSE	FALSE	FALSE	TRUE	FALSE
"3630_4"	FALSE	FALSE	FALSE	FALSE	FALSE
"9495_8"	FALSE	FALSE	FALSE	FALSE	TRUE
"8196_8"	FALSE	FALSE	TRUE	TRUE	TRUE

# FeatureHashing

- **FeatureHashing** implements **split** in the formula interface.

```
hash_size <- 2^16
m.train <- hashed.model.matrix(~ split(review, delim = " ", type = "existence"),
                               data = imdb.train,
                               hash.size = hash_size,
                               signed.hash = FALSE)
m.valid <- hashed.model.matrix(~ split(review, delim = " ", type = "existence"),
                               data = imdb.valid,
                               hash.size = hash_size,
                               signed.hash = FALSE)
```



## Type of **split**

- **existence**
- **count**
- **tf-idf** which is contributed by [Michaël Benesty](#) and will be announced in v0.9.1

# Gradient Boosted Decision Tree with **xgboost**

```
dtrain <- xgb.DMatrix(m.train, label = imdb.train$sentiment)
dvalid <- xgb.DMatrix(m.valid, label = imdb.valid$sentiment)
watch <- list(train = dtrain, valid = dvalid)
g <- xgb.train(booster = "gblinear", nrounds = 100, eta = 0.0001, max.depth = 2,
               data = dtrain, objective = "binary:logistic",
               watchlist = watch, eval_metric = "auc")
```

```
[0] train-auc:0.969895 valid-auc:0.914488
[1] train-auc:0.969982 valid-auc:0.914621
[2] train-auc:0.970069 valid-auc:0.914766
...
[97] train-auc:0.975616 valid-auc:0.922895
[98] train-auc:0.975658 valid-auc:0.922952
[99] train-auc:0.975700 valid-auc:0.923014
```

# Performance

- The AUC of **g** is 0.90110 in the public leader board
  - It outperforms the benchmark in Kaggle

# The Purpose of FeatureHashing

- Make our life easier



# Formula Interface in R

# Algorithm in Text Book

## Regression

$$y = X\beta + \varepsilon$$

## Real Data

	SEPAL.LENGTH	SEPAL.WIDTH	PETAL.LENGTH	PETAL.WIDTH	SPECIES
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
51	7.0	3.2	4.7	1.4	versicolor
52	6.4	3.2	4.5	1.5	versicolor
101	6.3	3.3	6.0	2.5	virginica
102	5.8	2.7	5.1	1.9	virginica

- How to convert the real data to  $X$ ?

# Feature Vectorization and Formula Interface

- $X$  is usually constructed via `model.matrix` in R

```
model.matrix(~ ., iris.demo)
```

	(INTERCEPT)	SEPAL.LENGTH	SEPAL.WIDTH	PETAL.LENGTH	PETAL.WIDTH	SPECIESVERSICOLOR	SPECIESVIRGINICA
1	1	5.1	3.5	1.4	0.2	0	0
2	1	4.9	3.0	1.4	0.2	0	0
51	1	7.0	3.2	4.7	1.4	1	0
52	1	6.4	3.2	4.5	1.5	1	0
101	1	6.3	3.3	6.0	2.5	0	1
102	1	5.8	2.7	5.1	1.9	0	1

# Formula Interface

- $y \sim a + b$ 
  - $y$  is the response
  - $a$  and  $b$  are predictors



# Formula Interface: +

- + is the operator of combining linear predictors

```
model.matrix(~ a + b, data.demo)
```

	(INTERCEPT)	A	B
1	1	5.1	3.5
2	1	4.9	3.0
51	1	7.0	3.2
52	1	6.4	3.2
101	1	6.3	3.3
102	1	5.8	2.7

# Formula Interface: :

- : is the interaction operator

```
model.matrix(~ a + b + a:b, data.demo)
```

	(INTERCEPT)	A	B	A:B
1	1	5.1	3.5	17.85
2	1	4.9	3.0	14.70
51	1	7.0	3.2	22.40
52	1	6.4	3.2	20.48
101	1	6.3	3.3	20.79
102	1	5.8	2.7	15.66

# Formula Interface: \*

- \* is the operator of cross product

```
# a + b + a:b  
model.matrix(~ a * b, data.demo)
```

	(INTERCEPT)	A	B	A:B
1	1	5.1	3.5	17.85
2	1	4.9	3.0	14.70
51	1	7.0	3.2	22.40
52	1	6.4	3.2	20.48
101	1	6.3	3.3	20.79
102	1	5.8	2.7	15.66

# Formula Interface: (

- $:$  and  $*$  are distributive over  $+$

```
# a:c + b:c
model.matrix(~ (a + b):c, data.demo)
```

	(INTERCEPT)	A:C	B:C
1	1	7.14	4.90
2	1	6.86	4.20
51	1	32.90	15.04
52	1	28.80	14.40
101	1	37.80	19.80
102	1	29.58	13.77



# Formula Interface: .

- . means all columns of the **data**.

```
# ~ Sepal.Length + Sepal.Width + Petal.Length +  
#   Petal.Width + Species  
model.matrix(~ ., iris.demo)
```

	(INTERCEPT)	SEPAL.LENGTH	SEPAL.WIDTH	PETAL.LENGTH	PETAL.WIDTH	SPECIESVERSICOLOR	SPECIESVIRGINIA
1	1	5.1	3.5	1.4	0.2	0	0
2	1	4.9	3.0	1.4	0.2	0	0
51	1	7.0	3.2	4.7	1.4	1	0
52	1	6.4	3.2	4.5	1.5	1	0
101	1	6.3	3.3	6.0	2.5	0	1
102	1	5.8	2.7	5.1	1.9	0	1

- Please check ?formula

# Categorical Features

# Categorical Feature in R

- A categorical variables of  $K$  categories are transformed to a  $K - 1$ -dimensional vector.
- There are many coding systems and the most commonly used is **Dummy Coding**.
  - The first category are transformed to  $\vec{0}$ .

## Dummy Coding

```
##          versicolor virginica
## setosa          0          0
## versicolor      1          0
## virginica        0          1
```

# Categorical Feature in Machine Learning

- Predictive analysis
- Regularization
- The categorical variables of  $K$  categories are transformed to  $K$ -dimensional vector.
  - The missing data are transformed to  $\vec{0}$ .

```
contr.treatment(levels(iris.demo$Species), contrasts = FALSE)
```

```
##          setosa versicolor virginica
## setosa         1          0          0
## versicolor     0          1          0
## virginica      0          0          1
```

# Motivation of the FeatureHashing



# Kaggle: Display Advertising Challenge

- Given a user and the page he is visiting, what is the probability that he will click on a given ad?



# The Size of the Dataset

- 13 integer features and 26 categorical features
- $7 \times 10^7$  instances
- Download the dataset via this [link](#)

## Vectorize These Features in R

- After binning the integer features, I got  $3 \times 10^7$  categories.
- Sparse matrix was required
  - A dense matrix required 14PB memory...
  - A sparse matrix required 40GB memory...

# Estimating Computing Resources

- Dense matrix: `nrow`  $(7 \times 10^7) \times$  `ncol`  $(3 \times 10^7) \times 8$  bytes
- Sparse matrix: `nrow`  $(7 \times 10^7) \times (13 + 26) \times (4 + 4 + 8)$  bytes

# Vectorize These Features in R

- `sparse.model.matrix` is similar to `model.matrix` but returns a sparse matrix.

# Fit Model with Limited Memory

- A post [Beat the benchmark with less than 200MB of memory](#) describes how to fit a model with limited resources.
  - online logistic regression
  - feature hash trick
  - adaptive learning rate

# Online Machine Learning

- In online learning, the model parameter is updated after the arrival of every new datapoint.
  - Only the aggregated information and the incoming datapoint are used.
- In batch learning, the model parameter is updated after access to the entire dataset.

# Memory Requirement of Online Learning

- Related to model parameters
  - Require little memory for large amount of instances.



# Limitation of `model.matrix`

- The vectorization requires **all categories**.

```
contr.treatment(levels(iris$Species))
```

```
##           versicolor virginica  
## setosa             0         0  
## versicolor        1         0  
## virginica         0         1
```

# My Work Around

- Scan all data to retrieve all categories
- Vectorize features in an online fashion
- The overhead of exploring features increases

# Observations

- Mapping features to  $\{0, 1, 2, \dots, K\}$  is one of method to vectorize feature.
  - **setosa**  $\Rightarrow \vec{e}_1$
  - **versicolor**  $\Rightarrow \vec{e}_2$
  - **virginica**  $\Rightarrow \vec{e}_3$

```
##          setosa versicolor virginica
## setosa          1           0         0
## versicolor      0           1         0
## virginica       0           0         1
```

# Observations


- `contr.treatment` ranks all categories to map the feature to integer.
- What if we do not know all categories?
  - Digital features are integer.
  - We use a function maps  $\mathbb{Z}$  to  $\{0, 1, 2, \dots, K\}$

```
charToRaw("setosa")
```

```
## [1] 73 65 74 6f 73 61
```

# What is Feature Hashing?

- A method to do feature vectorization with a hash function



- For example, Mod  $\% \%$  is a family of hash function.

# Feature Hashing

- Choose a hash function and use it to hash all the categorical features.
- The hash function does not require global information.

# An Example of Feature Hashing of Criteo's Data

V15	V16	V17
68fd1e64	80e26c9b	fb936136
68fd1e64	f0cf0024	6f67f7e5
287e684f	0a519c5c	02cf9876
68fd1e64	2c16a946	a9a87e68
8cf07265	ae46a29d	c81688bb
05db9164	6c9c9cf3	2730ec9c

- The categorical variables have been hashed onto 32 bits for anonymization purposes.
- Let us use the last 4 bits as the hash result, i.e. the hash function is `function(x) x %% 16`
- The size of the vector is  $(2^4)$ , which is called `hash size`

# An Example of Feature Hashing of Criteo's Data

V15	V16	V17
68fd1e64	80e26c9b	fb936136
68fd1e64	f0cf0024	6f67f7e5
287e684f	0a519c5c	02cf9876
68fd1e64	2c16a946	a9a87e68
8cf07265	ae46a29d	c81688bb
05db9164	6c9c9cf3	2730ec9c

- $68fd1e64, 80e26c9b, fb936136 \Rightarrow 4, b, 6$   
 -  $(0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0)$
- $68fd1e64, f0cf0024, 6f67f7e5 \Rightarrow 4, 4, 5$   
 -  $(0, 0, 0, 2, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0)$



# Hash Collision

- $68fd1e64, f0cf0024, 6f67f7e5 \Rightarrow 4, 4, 5$ 
  - $(0, 0, 0, 2, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0)$
- Hash function is a many to one mapping, so different features might be mapped to the same index. This is called collision.
- In the perspective of statistics, the collision in hash function makes the effect of these features confounding.

# Choosing Hash Function

- Less collision rate
  - Real features, such as text features, are not uniformly distributed in  $\mathbb{Z}$ .
- High throughput
- **FeatureHashing** uses the [Murmurhash3](#) algorithm implemented by [digest](#)

# Pros of FeatureHashing

## A Good Companion of Online Algorithm

```
library(FeatureHashing)
hash_size <- 2^16
w <- numeric(hash_size)
for(i in 1:1000) {
  data <- fread(paste0("criteo", i))
  X <- hashed.model.matrix(V1 ~ ., data, hash.size = hash_size)
  y <- data$V1
  update_w(w, X, y)
}
```

# Pros of FeatureHashing

## A Good Companion of Distributed Algorithm

```
library(pbdMPI)
library(FeatureHashing)
hash_size <- 2^16
w <- numeric(hash_size)
i <- comm.rank()
data <- fread(paste0("criteo", i))
X <- hashed.model.matrix(V1 ~ ., data, hash.size = hash_size)
y <- data$V1
# ...
```


# Pros of FeatureHashing

## Simple Training and Testing

```
library(FeatureHashing)
model <- is_click ~ ad * (url + ip)
m_train <- hashed.model.matrix(model, data_train, hash_size)
m_test <- hashed.model.matrix(model, data_test, hash_size)
```

# Cons of FeatureHashing

Hash Size



# Cons of FeatureHashing

## Lose Interpretation

- Collision makes the interpretation harder.
- It is inconvenient to reverse the indices to feature.

```
m <- hashed.model.matrix(~ Species, iris, hash.size = 2^4, create.mapping = TRUE)
hash.mapping(m) %% 2^4
```

```
##      Speciessetosa  Speciesvirginica Speciesversicolor
##                7                13                8
```

# The Result of the Competition...

- No.1: Field-aware Factorization Machine
  - The hashing trick do not contribute any improvement on the leaderboard. They apply the hashing trick only because it makes their life easier to generate features..
- No.3, no.4 and no.9 (we): Neuron Network



# Extending Formula Interface in R

- Formula is the most R style interface

## Tips

- `terms.formula` and its argument `specials`

# Specials

- `attr(tf, "specials")` tells which rows of `attr(tf, "factors")` need to be parsed further

```
rownames(attr(tf, "factors"))
```

```
## [1] "Plant"          "Type"           "conc"  
## [4] "split(Treatment)"
```

```
attr(tf, "specials")
```

```
## $split  
## [1] 4
```

# Parse

- **parse** extracts the information from the **specials**


```
options(keep.source = TRUE)
p <- parse(text = rownames(attr(tf, "factors"))[4])
getParseData(p)
```

##	line1	col1	line2	col2	id	parent	token	terminal	text
## 9	1	1	1	16	9	0	expr	FALSE	
## 1	1	1	1	5	1	3	SYMBOL_FUNCTION_CALL	TRUE	split
## 3	1	1	1	5	3	9	expr	FALSE	
## 2	1	6	1	6	2	9	' ('	TRUE	(
## 4	1	7	1	15	4	6	SYMBOL	TRUE	Treatment
## 6	1	7	1	15	6	9	expr	FALSE	
## 5	1	16	1	16	5	9	')'	TRUE	)

# Efficiency

- Rcpp
  - The core functions are implemented in C++
- Invoking external C functions
  - **digest** exposes the C function:
    - Register the c function
    - Add the helper header file
  - **FeatureHashing** imports the C function:
    - Import the c function

# Let's Discuss the Implementation Later



source: <http://www.commpraxis.com/wp-content/uploads/2015/01/commpraxis-audience-confusion.png>

# Summary

- Pros of **FeatureHashing**
  - Make it easier to do feature vectorization for predictive analysis.
  - Make it easier to tokenize the text data.
- Cons of **FeatureHashing**
  - Decrease the prediction accuracy if the **hash size** is too small
  - Interpretation becomes harder.

When should I use **FeatureHashing**?

Short Answer: Predictive Analysis with a large amount of categories.